

CS 433 Automated Reasoning 2025

Lecture 12: Satisfiability modulo theory (SMT) solvers

Instructor: Ashutosh Gupta

IITB India

Compile date: 2025-02-20

CDCL(\mathcal{T})

CDCL solves (i.e. checks satisfiability) quantifier-free propositional formulas

CDCL(\mathcal{T}) solves quantifier-free formulas in theory \mathcal{T} ,

- ▶ separates the boolean and theory reasoning,
- ▶ proceeds like CDCL, and
- ▶ needs support of a \mathcal{T} -solver $DP_{\mathcal{T}}$, i.e., a decision procedure for conjunction of literals of \mathcal{T}

The tools that are build using CDCL(\mathcal{T}) are called satisfiability modulo theory solvers (**SMT solvers**)

CDCL(\mathcal{T}) - some notation

Let \mathcal{T} be a first-order-logic theory with signature \mathbf{S} .

We assume input formulas are from \mathcal{T} , quantifier-free, and in CNF.

Definition 12.1

For a quantifier-free \mathcal{T} formula F , let $\text{atoms}(F)$ denote the set of atoms appearing in F .

Example 12.1

- ▶ $f(x) = g(h(x, y))$ is a formula in QF_EUF.
- ▶ $x > 0 \vee y + x = 3.5z$ is a formula in QF_LRA.

Boolean encoder

For a formula F , let **boolean encoder** e be a partial map from $atoms(F)$ to fresh boolean variables.

Definition 12.2

For a formula F , let $e(F)$ denote the term obtained by replacing each atom a by $e(a)$ if $e(a)$ is defined.

Example 12.2

Let $F = x < 2 \vee (y > 0 \vee x \geq 2)$ and $e = \{x < 2 \mapsto x_1, y > 0 \mapsto x_2\}$.
 $e(F) = x_1 \vee (x_2 \vee \neg x_1)$

Exercise 12.1

Consider boolean encoder $e = \{x < 2 \mapsto x_1, y > 0 \mapsto x_2\}$. Encode the following.

► $e(x \geq 2) =$

► $e(x + y \leq 0) =$

► $e(x < 2 \Rightarrow y \leq 0) =$

► $e(\top) =$

Partial model

Definition 12.3

For a boolean encoder e , a *partial model* m is an ordered partial map from $\text{range}(e)$ to \mathcal{B} .

Example 12.3

partial models $\{x \mapsto 0, y \mapsto 1\}$ and $\{y \mapsto 1, x \mapsto 0\}$ are not same.

CDCL(\mathcal{T}) will proceed by constructing partial models like CDCL.

Reverse encoder

Definition 12.4

For a partial model m of e , let $e^{-1}(m) \triangleq \{e^{-1}(x) \mid x \mapsto 1 \in m\} \cup \{\neg e^{-1}(x) \mid x \mapsto 0 \in m\}$

Example 12.4

Let $e = \{x < 2 \mapsto x_1, y > 0 \mapsto x_2\}$ and $m = \{x_1 \mapsto 0, x_2 \mapsto 1\}$.

$$e^{-1} = \{x_1 \mapsto x < 2, x_2 \mapsto y > 0\}$$

$$e^{-1}(m) = \{\neg(x < 2), y > 0\}$$

Exercise 12.2

Consider boolean encoder $e = \{x < 2 \mapsto x_1, y > 0 \mapsto x_2\}$. Encode the following.

► $e^{-1}(\{x_1 \mapsto 0\}) =$

► $e^{-1}(\{x_3 \mapsto 0\}) =$

► $e^{-1}(\{x_1 \mapsto 0, x_2 \mapsto 0\}) =$

► $e^{-1}(\emptyset) =$

Theory propagation

If we have partial assignment m , then we need to check if the theory accepts the assignment.

In other words, we need to know if $\bigwedge e^{-1}(m)$ is sat.

Example 12.5

In last example, we had $e^{-1}(m) = \{\neg(x < 2), y > 0\}$.

We ask if $\bigwedge e^{-1}(m) = \neg(x < 2) \wedge y > 0$ is sat. If no, we need to backtrack the assignments.

We assume that function `THEORYDEDUCTION` can check satisfiability of $\bigwedge e^{-1}(m)$.

CDCL(\mathcal{T})

Algorithm 12.1: CDCL(\mathcal{T})(formula G)

```
 $e := \text{CREATEENCODER}(G); F := e(G); m := \text{UNITPROPAGATION}(m, F); dl := 0; dstack := \lambda x.0;$   
do  
  // backtracking  
  while  $m \not\models F$  do  
    if  $dl = 0$  then return unsat;  
     $(C, dl) := \text{ANALYZECONFLICT}(m);$   
     $m.\text{resize}(dstack(dl)); F := F \cup \{C\}; m := \text{UNITPROPAGATION}(m, F);$   
  // Boolean decision  
  if  $F$  is unassigned under  $m$  then  
     $dstack(dl) := m.\text{size}(); dl := dl + 1; m := \text{DECIDE}(m, F); m := \text{UNITPROPAGATION}(m, F);$   
  // Theory propagation  
  if  $F$  is unassigned or sat under  $m$  then  
     $(Cs, dl') := \text{THEORYDEDUCTION}(\mathcal{T})(\bigwedge e^{-1}(m), m, dstack, dl);$  // Theory solving  
    if  $dl' < dl$  then  $\{dl = dl'; m.\text{resize}(dstack(dl));\};$   
     $F := F \cup e(Cs); m := \text{UNITPROPAGATION}(m, F);$   
while  $F$  is unassigned under  $m$  or  $m \not\models F$  or  $e^{-1}(m)$  is unsat;  
return sat
```

F is Boolean encoding of input G

Same as SAT solver CDCL

returns a clause set and a decision level

Topic 12.1

THEORYDEDUCTION

Theory propagation

THEORYDEDUCTION looks at the atoms assigned so far and checks

- ▶ if they are mutually unsatisfiable
- ▶ if not, are there other literals from G that are implied by the current assignment

Any implementation must comply with the following goals

- ▶ Correctness: boolean model is consistent with \mathcal{T}
- ▶ Termination: unsat partial models are never repeated

THEORYDEDUCTION

THEORYDEDUCTION solves conjunction of literals and returns a set of clauses and a decision level.

$$(Cs, dl') := \text{THEORYDEDUCTION}(\mathcal{T})(\bigwedge e^{-1}(m), m, dstack, dl)$$

Cs may contain the clauses of the form

$$(\bigwedge L) \Rightarrow \ell$$

where $\ell \in \text{lits}(F') \cup \{\perp\}$ and $L \subseteq e^{-1}(m)$.

Example : THEORYDEDUCTION

Example 12.6

If $\text{THEORYDEDUCTION}(\text{QF_LRA})(x > 1 \wedge x < 0, \dots)$ is called, the returned clauses will be

$$Cs := \{(x > 1 \wedge x < 0 \Rightarrow \perp)\}.$$

If $\text{THEORYDEDUCTION}(\text{QF_LRA})(x > 1 \wedge y > 0, \dots)$ is called, the returned clauses may be

$$Cs := \{(x > 1 \wedge y > 0 \Rightarrow x + y > 0), \dots\}.$$

Assuming $x + y > 0$
occurs in input

Specification of THEORYDEDUCTION

The output of THEORYDEDUCTION must satisfy the following conditions

- ▶ If $\bigwedge e^{-1}(m)$ is unsat in \mathcal{T} then Cs must contain a clause with $\ell = \perp$. dl' is the decision level immediately after which the unsatisfiability occurred (clearly stated shortly).
- ▶ if $\bigwedge e^{-1}(m)$ is sat then $dl' = dl$.

Example : CDCL(QF_EUF)

Example 12.7

Consider $F' = (x = y \vee y = z) \wedge (y \neq z \vee z = u) \wedge (z = x)$

$$e(F') = (x_1 \vee x_2) \wedge (\neg x_2 \vee x_3) \wedge x_4$$

After $F := e(F'); m := \text{UNITPROPAGATION}(m, F)$

$$m = \{x_4 \mapsto 1\}$$

After $m := \text{DECIDE}(m, F);$

$$m = \{x_4 \mapsto 1, x_2 \mapsto 0\}$$

After $m := \text{UNITPROPAGATION}(m, F)$

$$m = \{x_4 \mapsto 1, x_2 \mapsto 0, x_1 \mapsto 1\}$$

Example : CDCL(QF_EUF) II

Since $m = \{x_4 \mapsto 1, x_2 \mapsto 0, x_1 \mapsto 1\}$, $e^{-1}(m) = \{x = y, y \neq z, z = x\}$

After $(Cs, dl') := \text{THEORYDEDUCTION}(\text{QF_EUF})(x = y \wedge y \neq z \wedge z = x, ..)$

$Cs = \{x \neq y \vee y = z \vee z \neq x\}$, $dl' = 0$, $e(Cs) = \{\neg x_1 \vee x_2 \vee \neg x_4\}$

After $F := F \cup e(Cs)$; $m := \text{UNITPROPAGATION}()$

$m = \{x_4 \mapsto 1, x_2 \mapsto 0, x_1 \mapsto 1\} \leftarrow$ conflict with learned clause

Exercise 12.3

Complete the run

Theory propagation implementation - incremental solver

Theory propagation is implemented **using** incremental theory solvers.

Incremental solver $DP_{\mathcal{T}}$ for theory \mathcal{T}

- ▶ takes input constraints as a sequence of literals,
- ▶ has a data structure that defines the solver state and satisfiability of constraints seen so far.

Theory solver $DP_{\mathcal{T}}$ interface

A theory solver must provide the following interface.

- ▶ $\text{push}(\ell)$ - adds literal ℓ in “constraint store”
- ▶ $\text{pop}()$ - removes last pushed literal from the store
- ▶ $\text{checkSat}()$ - checks satisfiability of current store
- ▶ $\text{unsatCore}()$ - returns the set of literals that caused unsatisfiability

Definition 12.5

An *unsat core* of Σ is a subset (preferably minimal) of Σ that is unsat.

Commentary: We assume that push and pop call $\text{checkSat}()$ at the end of their execution. Therefore, explicit calls to $\text{checkSat}()$ are not necessary. However, practical tools allow users to choose the policy of calling $\text{checkSat}()$ - lazy vs. eager

Theory propagation implementation

Algorithm 12.2: THEORYDEDUCTION

Input: Set of literals Ls

Read only input: m partial model, $dstack$ decision depths, dl current decision level, input formula G

foreach $\ell \in Ls$ **do**

$DP_{\mathcal{T}}.push(\ell)$

if $DP_{\mathcal{T}}.checkSat() == unsat$ **then**

 // theory conflict

$Ls' := DP_{\mathcal{T}}.unsatCore(); dl' := \max\{dl'' \mid \exists \ell \in Ls', i. m[i] = e(\ell) \wedge dstack(dl'') < i\};$

return $(\neg \wedge Ls', dl')$

else

 //implied clauses

$Cs := \emptyset;$

foreach $\ell \in Lits(G)$ **do**

$DP_{\mathcal{T}}.push(\neg \ell);$

if $DP_{\mathcal{T}}.checkSat() == unsat$ **then**

$Ls' := DP_{\mathcal{T}}.unsatCore(); Cs := Cs \cup \{\neg \wedge Ls'\};$

$DP_{\mathcal{T}}.pop();$

return (Cs, dl)

dl' is the latest decision after which all literals in Ls' became true.

$Ls' = Ls$ will also be correct.
But, inefficient.

ℓ is called implied
literal and $\neg \ell \in Ls'$

Example: Theory deduction unsat example

Example 12.8

Consider $Ls = \{x = z, x = y, f(x) \neq f(y)\}$

First we will push all the literals to the theory solver.

$DP_{\mathcal{T}}.push(x = z); DP_{\mathcal{T}}.push(x = y); DP_{\mathcal{T}}.push(f(x) \neq f(y)).$

We will call $DP_{\mathcal{T}}.checkSat()$, which will return *unsat*.

We will call $DP_{\mathcal{T}}.unsatCore()$, which will return $\{x = y, f(x) \neq f(y)\}$.

The returned clause will be $x \neq y \vee f(x) = f(y)$.

Theory deduction will also return an appropriate decision level.

Example: Theory deduction sat example

Example 12.9

Consider $x = y \in Ls$ and assume $f(x) = f(y) \in Lits(G)$.

After pushing Ls , let us assume $DP_{\mathcal{T}}.checkSat()$ returns sat.

We search for implied clauses.

Since $f(x) = f(y) \in Lits(G)$, we will eventually call $DP_{\mathcal{T}}.push(f(x) \neq f(y))$.

We get unsatisfiability and unsat core, $\{x = y, f(x) \neq f(y)\}$.

We return $x \neq y \vee f(x) = f(y)$ among the implied clauses.

Topic 12.2

Example theory propagation implementation

Let us study implementation of DP_{EUF}

Decides conjunction of literals in the theory of EUF with interface
push, pop, checkSat, and unsatCore.

push, checkSat, and pop

► $DP_{EUF}.push$

Algorithm 12.3: $DP_{EUF}.push(t_1 \bowtie t_2)$

1 $IncrEUF(t_1 \bowtie t_2);$

► $DP_{EUF}.checkSat()$ { **return** *conflictFound*; }

► $DP_{EUF}.pop()$ is implemented by recording the time stamp of pushes and undoing all the mergers happened after the last push.

Exercise 12.4

Write pseudo code for $DP_{EUF}.pop()$

Unsat core

Algorithm 12.4: $DP_{EUF}.unsatCore()$

assume(*conflictFound* = 1);

Let $(t_1 \neq t_2)$ be the disequality that was violated;

return $\{t_1 \neq t_2\} \cup getReason(t_1, t_2)$;

Algorithm 12.5: $getReason(t_1, t_2)$

Let $(t'_1 = t'_2)$ be the merge operation that placed t_1 and t_2 in same class;


if $t'_1 = f(s_1, \dots, s_k) = f(u_1, \dots, u_k) = t'_2$ was derived due to congruence **then**

 | $reason := \bigcup_i getReason(s_i, u_i)$

else

 | $reason := \{t'_1 = t'_2\}$

return $getReason(t_1, t'_1) \cup reason \cup getReason(t'_2, t_2)$




Topic 12.3

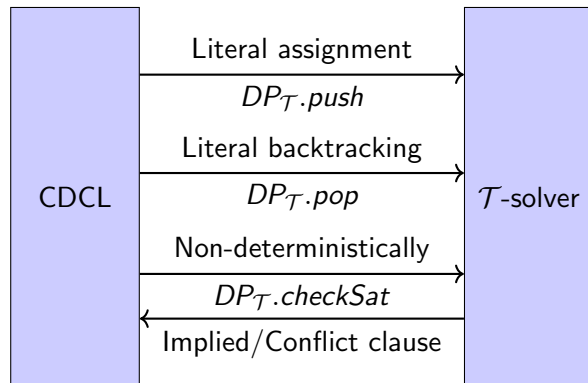
SMT Solvers

Incremental theory propagation

Earlier $CDCL(\mathcal{T})$



Fine-grained interaction with theory



Theory propagation strategies

- ▶ Exhaustive or Eager :
Cs contains all possible clauses
- ▶ Minimal or Lazy :
Cs only contains the clause that refutes current m
- ▶ Somewhat Lazy :
Cs contains only easy to deduce clauses

Rise of SMT solvers

- ▶ In early 2000s, stable SMT solvers started appearing. e.g., Yiecs
- ▶ SMT competition(SMT-comp) became a driving force in their ever increasing efficiency
- ▶ Formal methods community quickly realized their potential
- ▶ Z3, one of the leading SMT solver, alone has about 3000+ citations (375 per year)(June 2016)

Leading tools

The following are some of the leading SMT solvers

- ▶ Z3
- ▶ CVC4
- ▶ MathSAT
- ▶ Boolector

Topic 12.4

Problems

Run SMT solvers

Exercise 12.5

- Find a satisfying assignment of the following formula using SMT solver

$$(x > 0 \vee y < 0) \wedge (x + y > 0 \vee x - y < 0)$$

Give the model generated by the SMT solver.

- Prove the following formula is valid using SMT solver

$$(x > y \wedge y > z) \Rightarrow x > z$$

Give the proof generated by the SMT solver.

Please do not simply submit the output. Please write the answers in the mathematical notation.

Knapsack problem

Exercise 12.6

Write a program for solving the knapsack problem that requires filling a knapsack with stuff with maximum value. For more information look at the following.

`https://en.wikipedia.org/wiki/Knapsack_problem`

The output of the program should be the number of solutions that have value more than 95% of the best value.

Download Z3 from the following webpage: `https://github.com/Z3Prover/z3`

We need a tool to feed random inputs to your tool. Write a tool that generates random instances, similar to what was provided last time.

Evaluate the performance on reasonably sized problems. You also need to design the evaluation strategy. Evaluation plots and a small text to describe your strategies.

Topic 12.5

Extra slides : optimizations

Implied literals without implied clauses

Bottleneck: There may be too many implied clauses.

Observation: Very few of the implied clauses are useful, i.e., contribute in early detection of conflict.

Optimization: apply implied literals, without adding implied clauses.

Optimization overhead: If an implied model is used in conflict then recompute the implied clause for the implication graph analysis.

Relevancy

Bottleneck: All the assigned literals are sent to the theory solver.

Observation: However, *CDCL* only needs to send those literals to the solver that make unique clauses satisfiable.

Optimization:

- ▶ Each clause chooses one literal that makes it sat under current model.
- ▶ Those clause that are not sat under current model do nothing.
- ▶ If a literal is not chosen by any clause then it is not passed on to \mathcal{T} -solver.

Patented: US8140459 by Z3 guys (the original idea is more general than stated here)

Optimization overhead: Relevant literal management

Exercise 12.7

Suggest a scheme for relevant literal management.

End of Lecture 12